

Prognos: Solution Document

Connected Vehicle Intelligence Hackathon

To be submitted by	Pratibimb Gupta (RA2311003010027). Team name: Prognos
Team members and roles	Pratibimb Gupta: sole member, all roles. Email: codingupta@gmail.com
Problem space chosen	Predictive maintenance, with a fleet-wide emerging-fault radar
Repository URL	https://github.com/championorwhat/Connected-Vehicle-Intelligence
Explainer video	In folder 2 - Hackathon Explainer Video of the submission Google Drive folder
Date of submission	02/10/2026

How to read the numbers. Every measured number links to the evidence file that produced it. Anything not measured says **NOT YET MEASURED**. Unless stated otherwise, measurements come from a **4-vCPU x86_64 Linux container shared by every service**. That is not the author's target machine (an 8 GB MacBook Air M2), on which nothing has been measured yet. All data is synthetic.

Table of contents

1. [Executive Summary](#)
 2. [Problem Statement & Validation](#)
 3. [Solution Description](#)
 4. [Feature List](#)
 5. [Solution Architecture \(High-Level Design\)](#)
 6. [Low-Level Design](#)
 7. [Non-Functional Requirements & Performance Benchmarks](#)
 8. [Security & Compliance](#)
 9. [Test Strategy](#)
 10. [Observability](#)
 11. [AI / ML Component](#)
 12. [Architecture Decisions, Risks & Future Enhancements](#)
 13. [Demo Video](#)
 14. [Repository Checklist](#)
 15. [Conclusion](#)
 16. [Declarations](#)
 17. [Appendix](#)
-

1. Executive Summary

Problem. A fleet maintenance manager with 100,000 mixed petrol, diesel and electric vehicles learns about most faults when a vehicle breaks down at the roadside or a warning light comes on. At that scale nobody can triage the signals by hand, so repairs are unplanned, and unplanned repairs cost more than planned ones.

Solution. Prognos reads telemetry from vehicles of three different makes, raises a critical alert within seconds, predicts which vehicles are likely to fail within 7 days and why, and proposes a workshop booking before the predicted failure.

Key results (measured; 4-vCPU Linux container):

What	Result	Target
100K vehicles, one event every 10 s each (10K events/s)	every pipeline stage keeps up (evidence)	100K vehicles at ~1 event/s each
Critical alert latency, 100K vehicles	p95 1.78 s (evidence)	< 5 s
Event → dashboard, 100K vehicles	p95 1.88 s (evidence)	< 2 s
API with 40 users during the 100K run	p95 63–82 ms, 0 errors (evidence)	p95 < 200 ms
Breakdowns warned before they happened (simulated back-test)	98.6 % , median 49.6 h ahead (evidence)	–
Failure model vs rules, held-out data	PR-AUC 0.742 vs 0.508 (evidence)	beat the baseline
End-to-end ingest at 100K events/s	not reached : this 4-vCPU machine sustains about 15K events/s end to end (evidence)	100K+ events/s

What is different.

- **Evaluated against ground truth.** The simulator records when each injected fault really fails, so recall, lead time and model quality are measured, not asserted.
- **Honest money.** Work orders are ranked by calibrated risk and capacity. The system refuses to show a money figure until repair costs have a cited source.
- **Emerging-fault radar.** It finds a fault code that suddenly spikes in one firmware release, using exact statistics, with no false signals in the tests.
- **Measured, then fixed.** Load tests found a crash loop, an API freeze and a disk-filling Kafka setting. SQL tuning cut the replayed dashboard workload from 195.6 s to 23.8 s of API time. Each fix has evidence.

2. Problem Statement & Validation

2.1 Problem Statement

A fleet maintenance manager running a 100,000-vehicle mixed ICE/EV fleet **needs a way to** know which vehicles are likely to fail in the next 7 days, and why, early enough to book them into a workshop, **because** faults today surface as roadside breakdowns or warning lights that nobody can triage across tens of thousands of vehicles, **which today costs** [unplanned breakdowns per 1,000 vehicles per month] × [(unplanned – planned repair cost) + downtime cost per day × extra downtime days].

The cost is a **formula, not a number**. Its parameters are **NOT SOURCED** (see 2.2), and Prognos does not invent them.

Primary user: the fleet maintenance manager, who owns vehicle uptime and the workshop schedule.

Secondary stakeholders:

Stakeholder	Interest
Workshop technicians	A work order with the suspected component and the evidence, not just a fault code
Drivers	Fewer roadside breakdowns; their location is minimised and masked
Fleet finance	Cost avoided against inspection cost (needs sourced costs)
OEM quality and warranty teams	A fault spreading across one model or firmware release
Data protection officer	Retention, erasure and an audit trail (GDPR; India's DPDP Act 2023)

The choice among the eight official problem spaces is scored in the [M0 decision matrix](#) (section 2). Predictive maintenance won because every 1 Hz signal matters to it, it needs both the real-time and the batch path, and a simulator can supply ground truth.

2.2 Evidence & Validation

Facts vs assumptions.

Evidence / Assumption	Source or method	What it shows	Confidence
Fact (brief): 100K vehicles, 100K+ events/s, critical alerts < 5 s, dashboard < 2 s	Hackathon problem statement	The scale and latency the product must meet	High
Fact (industry): US truck repair and maintenance averaged 21.5 US cents per mile in 2025, up 8.6 % on 2024 (parts, labour and roadside service; excluding tyres and towing)	ATRI, <i>An Analysis of the Operational Costs of Trucking</i> (2026 update), as reported by trade press (FleetOwner , Transport Topics). The original report was not read (not reachable from the build environment)	Maintenance is a large, rising line item	Medium (secondary reporting)
Assumption: an unplanned roadside repair costs several times a planned repair of the same component (towing, emergency labour, downtime)	Vendor blogs claim 3–9×; no independent study was read	Why finding faults early pays	Low. Used in no calculation
Fact (simulated): trend rules warn of 98.6 % of breakdowns, a median of 49.6 h ahead	Back-test against the simulator's ground truth (evidence)	Early warning is possible when faults degrade gradually	High for the simulator; real fleets NOT YET MEASURED

Evidence / Assumption	Source or method	What it shows	Confidence
Fact (simulated): a gradient-boosted model beats the rules on held-out runs (PR-AUC +0.23, 95 % interval [0.17, 0.30])	Held-out runs and two shifted variants (evidence)	Learning from telemetry adds value over fixed rules	Medium: may learn the simulator
Assumption: a mixed fleet uses several telematics sources with different formats	Design premise; modelled as 3 OEM formats	Normalisation is needed before any analysis	Medium

Validation method. Simulation with ground truth. The simulator ([algorithms A1-A4](#)) degrades five components along plausible curves and records the true failure time. Detection, calibration and the model are all scored against it, on runs they were not tuned on. No interviews were held, and no public fleet dataset with failure labels was used.

Existing alternatives, and why they fall short for this user.

Alternative	Falls short because
OEM portals	Each covers one brand; a mixed fleet gets several portals and no single ranked list
Aftermarket dongles and generic telematics	Report fault codes and positions; the triage of tens of thousands of codes is left to the manager
Integration platforms (e.g. Motorq Fuse)	Normalise and deliver vehicle data; deciding which vehicle to book, and where, is left to the customer. This is our reading of the public positioning; no product was tested
Fixed-interval servicing	Ignores the vehicle's condition: services healthy vehicles and misses fast faults

2.3 Impact & Success Metrics

Metric	Baseline today	Target	Achieved	How measured
Breakdowns warned ≥ 48 h ahead (recall, top-5 % capacity list)	Rules: 23.2 %	Beat the rules	Model: 34.7 %	Held-out back-test (evidence)
Precision of the at-risk list (top 5 %)	Rules: 72.5 %	Beat the rules	Model: 99.7 %	Same back-test
Median prediction lead time	Rules: 38.1 h	≥ 72 h	Model: 50.4 h (target not met)	Same back-test
Expected cost avoided per 1,000 vehicles per month	-	Positive after inspection cost	NOT COMPUTED: costs not sourced	Back-test $\times$ sourced costs
Critical alert latency	-	p95 < 5 s	p95 1.78 s at 100K vehicles	Live load test (evidence)
Event $\rightarrow$ dashboard	-	< 2 s	p95 1.88 s at 100K vehicles	WebSocket client (evidence)

Scale of impact. The pipeline's cost grows with events, not with users. On the same 4 vCPUs, 10K vehicles used 0.9 CPUs and 1.5 GiB; 100K vehicles used 2.9 CPUs and 3.1 GiB ([load tests](#)). The manager's work does not grow that way: the at-risk list is capped at workshop capacity, so a 100K fleet produces a bounded daily list of proposed bookings rather than 10 $\times$ more alerts to read.

Wider impact. Safety: fewer vehicles failing in traffic (tyre and overheating faults are raised within seconds). Environment: fewer tow-truck trips, and a planned repair replaces only the failing part. Compliance: location is minimised, masked and erasable (§8). These are qualitative; none is measured.

3. Solution Description

3.1 Solution Overview & User Journey

What it does. Prognos collects telemetry from three vehicle makes, converts it into one format, watches every vehicle continuously, and turns what it sees into three things for the fleet manager: an immediate alert when something is critical, a ranked list of vehicles likely to fail this week with the reasons, and proposed workshop bookings.

User journey.

Step	What happens	Where
1. Vehicle	Emits telemetry in its maker's format; a degrading part shifts its signals	apps/simulator
2. Ingest	Kafka <code>telemetry.raw</code> , keyed by vehicle	ADR-001
3. Normalise	Map to the canonical event, validate, drop duplicates; rejects go to the DLQ with a reason	normalizer
4. Detect	Critical rules fire at once; trend statistics raise early warnings	detector
5. Store	Alerts → PostgreSQL; live state → Redis; history → ClickHouse	sink , ClickHouse Kafka engine
6. Predict	Every 5 minutes the model scores the fleet (shadow mode)	scorer
7. Plan	Open alerts become proposed work orders at the nearest workshop with capacity	planner
8. Act	The manager sees the alert live, acknowledges it and schedules the work order	dashboard
9. Outcome	The technician completes the order with an outcome; every step is audited	API, audit_log

Screenshots of the working product (more in the README):

2,123

Vehicles in your fleet

7

Critical alerts open need action now

45

Warnings open each sign of a fault

16

Work orders proposed bookings waiting for you

0

Scheduled booked into a workshop

0

In workshop being repaired

Most at risk


Health score from 0 to 100: 100 means no problems, lower means more urgent. Click a vehicle for details.

Vehicle	Model	Health	What is wrong	Last report
PG2PK5C82NC000001	Vega Pickup	30 Needs attention	Misfire in several cylinders Misfire in cylinder 1 Engine running rough (possible misfire)	2s ago
PG1SD3H16AC000779	Orion Sedan Hybrid	30 Needs attention	Electrical system voltage low 12 V battery critically low 12 V battery getting weak	2s ago
PG1HL2873TC000695	Orion Hauler	30 Needs attention	Low tyre pressure, rear left Tyre pressure critically low Tyre slowly losing pressure	10s ago
PG3ET9T513LC000269	Lyra e-Truck	30 Needs attention	Low tyre pressure, front right Tyre pressure critically low Tyre slowly losing pressure	4s ago
PG3EC4S64NC000003	Lyra e-Compact	30 Needs attention	Low tyre pressure, rear right Tyre pressure critically low Tyre slowly losing pressure	2s ago

Where they are

The vehicles in the list. Red: needs attention; amber: keep an eye on. Click a dot to open it.

<



The screenshots were taken without internet access, so the map tiles are blank.

3.2 Key Value Proposition

- **Customer job:** keep the fleet on the road by repairing vehicles before they break down, within the workshops' capacity.
- **Pain relieved:** fewer roadside breakdowns, and a short, ranked list in place of thousands of raw fault codes. Critical faults arrive within seconds.
- **Gain created:** a prediction for each vehicle (likely to fail within 7 days, and why), a proposed booking before the failure, and early notice when a fault spreads across a model or firmware release.
- **Differentiation:** Prognos is judged against ground truth, not demos. It says "not sourced" rather than invent a saving, keeps the model in shadow until it is validated on live outcomes, and measures its own latency against published SLOs.

3.3 Innovative Ideas

Idea	What is new	Evidence
1. A simulator with ground truth, used as the test oracle	Telemetry demos usually cannot say whether a prediction was right. Here five failure modes degrade along physical signatures, and the simulator publishes the true failure time. Rule recall, calibration, deadlines and the ML model are scored against it on runs they were not fitted on.	Back-test: 98.6 % of breakdowns warned, median lead 49.6 h (evidence); calibration held-out on another seed, 23 of 24 rules inside the 95 % interval (evidence)
2. A capacity-aware planner that refuses fabricated money	Each alert rule gets a measured probability and a conservative deadline (10th-percentile time to failure). The planner books the most valuable vehicles first into the nearest workshop with capacity before that deadline. Value in money is computed only from costs that cite a source; the database rejects a cost without one.	100K vehicles, 5,000 open alerts: a cycle takes 0.90 s, and a re-run proposes nothing new in 28 ms (evidence); test_cost_parameters_require_provenance
3. Emerging-fault radar with exact statistics	Compares each model or firmware cohort with like-for-like peers, counting distinct vehicles, with an exact Poisson tail and a Bonferroni correction. The M0 plan proposed a Count-Min Sketch; it was rejected because the keys are few and a sketch only adds error.	Defect found in every window, 0 false signals on a control run (evidence); live at 100K vehicles, only the defective cohort flagged (evidence)

4. Feature List

Every feature is traceable to code, a test and evidence in [docs/feature-traceability.csv](#), which has the full columns. The table below is generated from that file (`scripts/sync_solution_document.py`); a unit test fails if they differ. Video times marked "(script)" come from the [demo script](#); they are confirmed once the video is recorded.

ID	Feature	User story	Priority	Status	Code path	Video
F-00	One-command local stack	As a developer I want one command to start the data plane so that anyone can run Prognos locally	Must	Done	<code>docker-compose.yml</code> <code>Makefile</code>	-
F-00b	Replicated durable messaging	As a platform engineer I want telemetry to survive a broker failure so that no events are lost	Must	Partial	<code>infra/docker/compose.kafka-ha.yml</code> <code>kafka/topics/</code>	-
F-01	Multi-tenant relational core (3NF)	As a fleet manager I want my fleet data isolated from other customers so that it is never exposed or corrupted	Must	Done	<code>database/postgres/migrations/</code>	-
F-02	Telemetry store with duplicate-safe rollups	As a fleet manager I want vehicle history kept for trend analysis so that slow degradation is visible	Must	Done	<code>database/telemetry/migrations/</code>	-
F-03	Deterministic 100K-vehicle synthetic fleet	As a platform engineer I want a reproducible 100K-vehicle dataset so that every benchmark is repeatable	Must	Done	<code>packages/common/src/prognos_common/roster.py</code> <code>database/postgres/seeds/load_seed.py</code>	-
F-04	VIN and DTC validation	As a platform engineer I want malformed identifiers rejected so that bad OEM data never reaches analytics	Must	Done	<code>packages/common/src/prognos_common/vin.py</code> <code>packages/common/src/prognos_common/dtc.py</code>	-
F-05	Append-only audit log	As a DPO I want an immutable record of data access so that accountability can be demonstrated	Must	Partial	<code>database/postgres/migrations/20260929000001_core_schema.sql</code>	-
F-06	Scalable multi-OEM telemetry simulator	As a platform engineer I want realistic 100K-vehicle telemetry in three OEM formats so that the pipeline is tested against real-world variety and scale	Must	Done	<code>apps/simulator/src/prognos_sim/</code>	-
F-07	Fault injection with ground truth	As a data scientist I want known failure onset and failure times so that predictions can be evaluated honestly	Must	Done	<code>apps/simulator/src/prognos_sim/fleet.py</code>	-
F-08	Delivery anomaly injection (duplicates / out-of-order / malformed / missing)	As a platform engineer I want realistic bad data so that validation, dedup and DLQ are proven	Must	Done	<code>apps/simulator/src/prognos_sim/engine.py</code> <code>apps/simulator/src/prognos_sim/formats.py</code>	-

ID	Feature	User story	Priority	Status	Code path	Video
F-09	Multi-OEM normalisation to a canonical event	As a platform engineer I want every OEM format mapped to one schema so that downstream logic is written once	Must	Done	apps/stream-processor/src/prognos_stream/adapters.py	-
F-10	Validation with dead-letter queue	As a platform engineer I want invalid payloads quarantined with a reason so that bad data never reaches analytics and can be replayed	Must	Done	apps/stream-processor/src/prognos_stream/processor.py apps/stream-processor/src/prognos_stream/service.py	-
F-11	Exact duplicate suppression with out-of-order tolerance	As a fleet manager I want each event counted once so that alerts and costs are not double counted	Must	Done	apps/stream-processor/src/prognos_stream/dedup.py	-
F-12	Critical real-time alerts (< 5 s)	As a fleet manager I want to know within seconds when a vehicle overheats or breaks down so that I can act before damage spreads	Must	Done	apps/stream-processor/src/prognos_stream/detector.py	01:15 (script)
F-13	Early-warning trend detection	As a fleet manager I want warnings days before a breakdown so that repairs can be planned	Must	Done	apps/stream-processor/src/prognos_stream/features.py apps/stream-processor/src/prognos_stream/detector.py	01:40 (script)
F-14	Tyre leak rate and time-to-critical estimate	As a fleet manager I want to know how long a leaking tyre can keep running so that I can schedule the fix	Should	Done	apps/stream-processor/src/prognos_stream/detector.py	-
F-15	Detection back-test harness	As a data scientist I want to score detection against ground truth so that every change to rules or models is measured	Must	Done	apps/stream-processor/src/prognos_stream/evaluate.py	04:15 (script)
F-16	Telemetry history in ClickHouse via Kafka engine	As an analyst I want every event stored for trend analysis so that history is complete	Must	Done	database/telemetry/migrations/20260930000002_kafka_ingestion.sql	-
F-17	Idempotent alert persistence	As a fleet manager I want each alert recorded once with its lifecycle so that I can act on and audit it	Must	Done	apps/stream-processor/src/prognos_stream/sinks.py	-
F-18	Live vehicle state and at-risk ranking in Redis	As a fleet manager I want a live view ranked by risk so that I see the worst vehicles first	Must	Done	apps/stream-processor/src/prognos_stream/sinks.py	-
F-19	End-to-end no-data-loss reconciliation	As a platform engineer I want proof that nothing is lost between Kafka and the stores so that the brief's no-loss requirement is evidenced	Must	Done	scripts/reconcile.py	-

ID	Feature	User story	Priority	Status	Code path	Video
F-20	Calibrated failure risk per rule	As a fleet manager I want each warning to say how likely and how soon a failure is so that I can compare vehicles	Must	Done	apps/stream-processor/src/prognos_stream/evaluate.py scripts/build_calibration.py apps/stream-processor/src/prognos_stream/calibration/rules-v1.json	-
F-21	Cost-aware prioritisation (no fabricated money)	As a finance owner I want vehicles ranked by expected cost avoided only when the costs are real so that no decision rests on invented figures	Must	Done	apps/stream-processor/src/prognos_stream/planner.py	02:05 (script)
F-22	Capacity-aware work-order proposals with audit trail	As a workshop planner I want proposed bookings that respect daily capacity and the failure deadline so that the most valuable repairs happen first	Must	Done	apps/stream-processor/src/prognos_stream/planner.py apps/stream-processor/src/prognos_stream/planner_main.py	02:05 (script)
F-23	Emerging-fault radar (firmware / model cohorts)	As an OEM quality engineer I want to know when one firmware or model shows a DTC far more often than its peers so that a systematic defect is caught before it spreads	Should	Done	apps/stream-processor/src/prognos_stream/radar.py apps/stream-processor/src/prognos_stream/radar_main.py apps/simulator/src/prognos_sim/engine.py (firmware_defect)	02:35 (script)
F-24	Backfill training dataset from the production pipeline	As a data scientist I want training data produced by the same normalizer and detector as the live system so that the model learns what it will see in production	Must	Done	ml/src/prognos_ml/dataset.py	-
F-25	Leak-free window features (DuckDB over Parquet)	As a data scientist I want features that only use the past so that offline results hold in production	Must	Done	ml/src/prognos_ml/features.py	-
F-26	7-day failure model vs rule baseline on held-out data	As a fleet manager I want a better at-risk list than the rules so that workshop slots go to vehicles that will really fail	Must	Done	ml/src/prognos_ml/model.py ml/src/prognos_ml/pipeline.py	04:15 (script)
F-27	Per-prediction explanations and versioned model artefact	As a fleet manager I want to see why a vehicle is flagged so that I trust and can challenge the list	Must	Done	ml/src/prognos_ml/model.py ml/models/failure-7d-v1/	-

ID	Feature	User story	Priority	Status	Code path	Video
F-28	Secure login and tokens (RS256 JWT + JWKS; argon2id; brute-force guard)	As a fleet manager I want a secure login so that only my staff see my fleet	Must	Done	apps/api/src/prognos_api/security.py apps/api/src/prognos_api/routers/auth.py	-
F-29	Role-based access and tenant isolation with audit of denials	As a data owner I want each role to see only what it needs and no tenant to see another so that data is protected	Must	Done	apps/api/src/prognos_api/deps.py apps/api/src/prognos_api/routers/	-
F-30	Fleet REST API (vehicles; alerts; work-order lifecycle) with keyset pagination	As a dashboard developer I want stable paginated endpoints so that the UI stays fast at 100K vehicles	Must	Done	apps/api/src/prognos_api/routers/	-
F-31	Live alert feed over WebSocket	As a fleet manager I want new alerts pushed to my screen so that I react within seconds	Must	Done	apps/api/src/prognos_api/routers/live.py	01:15 (script)
F-32	Live model scoring with training/serving parity and a data gate	As a data scientist I want production scores computed exactly like training features so that offline results carry over	Must	Done	ml/src/prognos_ml/scorer.py ml/src/prognos_ml/features.py	-
F-33	Shadow deployment of the model (rules stay in charge until validated)	As a fleet manager I want the model introduced safely so that an unvalidated model never drives bookings	Must	Done	apps/stream-processor/src/prognos_stream/planner_main.py apps/api/src/prognos_api/routers/vehicles.py	01:00 (script)
F-34	Fleet dashboard: summary; at-risk list with rules/model toggle; map	As a fleet manager I want one screen that shows which vehicles need attention and where so that I can act in seconds	Must	Done	apps/web/src/pages/Overview.tsx apps/web/src/components/FleetMap.tsx	01:00 (script)
F-35	Live alerts view with acknowledge	As a fleet manager I want new alerts to appear without refreshing so that nothing is missed	Must	Done	apps/web/src/pages/Alerts.tsx	01:15 (script)
F-36	Work-order board with role-aware lifecycle actions	As a workshop planner I want to schedule, start and complete work orders so that repairs are tracked to an outcome	Must	Done	apps/web/src/pages/WorkOrders.tsx	02:05 (script)
F-37	Vehicle detail with plain-language model reasons	As a fleet manager I want to know why a vehicle is flagged so that I can trust or challenge it	Should	Done	apps/web/src/pages/VehicleDetail.tsx apps/web/src/format.ts	01:40 (script)
F-38	SLOs with multi-window burn-rate alerts and runbooks	As an operator I want to be paged only when users are affected so that alerts are trusted	Must	Done	infra/monitoring/prometheus/rules/prognos.rules.yml docs/observability/runbooks.md	-

ID	Feature	User story	Priority	Status	Code path	Video
F-39	Service-health dashboard generated from code	As an operator I want one screen that shows which pipeline stage is failing so that I can act fast	Must	Done	scripts/gen_dashboards.py infra/monitoring/grafana/dashboards/prognos-overview.json	03:00 (script)
F-40	Structured JSON logs with request ids	As an operator I want to follow one request through the logs and audit log so that I can explain a failure	Should	Done	packages/common/src/prognos_common/logs.py apps/api/src/prognos_api/main.py	-
F-41	Failure drill: detect a lost service and recover from the backlog	As an operator I want proof that outages are detected so that I can rely on the monitoring	Should	Done	infra/monitoring/prometheus/rules/prognos.rules.yml	03:25 (script)
F-42	PostgreSQL row-level security as a second tenant barrier	As a fleet customer I want my data isolated even if the application has a bug so that competitors never see it	Must	Done	database/postgres/migrations/20261002000004_row_level_security.sql apps/api/src/prognos_api/deps.py	-
F-43	Right-to-erasure workflow (DPO request + worker across PostgreSQL / ClickHouse / Redis)	As a data protection officer I want to erase a driver's or vehicle's personal data and prove it so that the fleet meets privacy law	Must	Done	apps/api/src/prognos_api/routers/privacy.py apps/api/src/prognos_api/erase.py	-
F-44	STRIDE threat model with controls mapped to tests	As a security reviewer I want every threat tied to a control and a test so that claims can be checked	Must	Done	docs/security/threat-model.md	-
F-45	Security headers and loopback-only ports	As an operator I want safe defaults so that a laptop demo does not expose data	Must	Done	apps/web/security-headers.inc docker-compose.yml	-
F-46	Dependency and secret scanning gate	As a maintainer I want vulnerable dependencies and leaked secrets to fail the build so that they never ship	Should	Done	.github/workflows/security.yml .trivyignore .github/dependabot.yml	-
F-47	Measured SQL optimisation of the three slowest queries	As a fleet manager I want the dashboard to load instantly even with years of alert history so that I can act without waiting	Must	Done	apps/api/src/prognos_api/routers/live.py apps/api/src/prognos_api/routers/work_orders.py database/postgres/migrations/20261003000005_alerts_active_index.sql	-


ID	Feature	User story	Priority	Status	Code path	Video
F-48	Combined coverage gate (unit + integration >= 80 %)	As a maintainer I want untested code to fail the build so that quality does not erode	Must	Done	<code>.github/workflows/ci.yml</code> <code>pyproject.toml</code> <code>Makefile</code>	-
F-49	BDD acceptance scenarios for the core user journeys	As a product owner I want the main journeys written as executable scenarios so that behaviour is checked in plain language	Should	Done	<code>tests/integration/features/fleet_manager.feature</code>	-
F-50	Chaos drill: database outage without data loss	As an operator I want proof that a database outage loses nothing so that I can trust recovery	Should	Done	<code>apps/stream-processor/src/prognos_stream/sink_main.py</code>	03:25 (script)
F-51	Container image vulnerability gate	As a security reviewer I want shipped images scanned so that known critical flaws never deploy	Should	Done	<code>.github/workflows/ci.yml</code> <code>.github/dependabot.yml</code>	-
F-52	Load tests at 10K/50K/100K vehicles with capacity ramp and burst	As an operator I want to know how many vehicles the platform handles and how it degrades so that I can size it	Must	Done	<code>scripts/load_test.py</code> <code>scripts/api_load.py</code>	03:25 (script)
F-53	Detector memory sized for 100K vehicles (crash-loop fix)	As an operator I want the detector to fit its memory limit at full fleet size so that alerts are not delayed by restarts	Must	Done	<code>apps/stream-processor/src/prognos_stream/detector.py</code>	-
F-54	Login bursts do not stall the API	As a fleet manager I want the dashboard to stay fast when everyone signs in at shift start so that I can work immediately	Must	Done	<code>apps/api/src/prognos_api/routers/auth.py</code> <code>apps/api/src/prognos_api/deps.py</code>	-
F-55	Crash-loop alert (ServiceRestarting)	As an operator I want to be paged when a service keeps restarting so that silent degradation is caught	Should	Done	<code>infra/monitoring/prometheus/rules/prognos.rules.yml</code>	-
F-56	Dashboard freshness measured (event to WebSocket < 2 s)	As a fleet manager I want a critical alert on my screen within 2 seconds so that I can act while the vehicle is still on the road	Must	Done	<code>scripts/ws_lateness.py</code> <code>apps/api/src/prognos_api/routers/live.py</code>	01:15 (script)
F-57	Open-source declaration from image SBOMs	As a reviewer I want every shipped component and its licence listed so that the submission's licence obligations are clear	Must	Done	<code>scripts/gen_licences.py</code> <code>evidence/sbom/</code>	-
F-58	Solution Document with evidence links checked in CI	As a judge I want every claim in the Solution Document to link to working evidence so that I can verify it	Must	Done	<code>docs/solution-document/solution-document.md</code> <code>scripts/sync_solution_document.py</code>	-

60 features: 58 Done, 2 Partial or Planned.

5. Solution Architecture (High-Level Design)

5.1 Architecture Overview


System context (C4 level 1).



There is no external identity provider, LLM API or data warehouse. Tokens are RS256 JWTs published as JWKS, so an OIDC provider can replace the built-in issuer through configuration.

Containers (C4 level 2). Every service is a container in [docker-compose.yml](#).

- **Arrows between services are Kafka topics** (Kafka 4.1, KRaft). Rejected messages go to `telemetry.dlq`.
- **Arrows to a store name the protocol**: SQL over the PostgreSQL wire protocol, ClickHouse over HTTP, Redis over RESP.
- **Left out for readability**: the `erasure-worker` (SQL and RESP to all three stores) and Prometheus, which scrapes `/metrics` over HTTP from every service.



One telemetry event, from vehicle to insight (100K vehicles, 10K events/s; measured latencies):

Hop	What happens	Measured latency
Vehicle → <code>telemetry.raw</code> → normalizer → <code>telemetry.canonical</code>	Includes the simulator's injected network delay (mean 0.3 s)	p50 0.35 s, p95 1.04 s (on-time events, evidence)
Event → alert opened by the detector	Rules and trend statistics	p50 0.62 s, p95 1.78 s (evidence)
Alert detected → dashboard WebSocket	<code>alerts</code> → sink → PostgreSQL + Redis pub/sub → API → socket	p50 0.23 s, p95 1.03 s (evidence)
Event → dashboard, end to end		p50 0.82 s, p95 1.88 s (same evidence)
Event → ClickHouse history	Kafka engine, flushed in blocks	not measured per event
Telemetry → model score	Scorer cycle over 100K vehicles	about 17 s per cycle (evidence)

5.2 Technology Stack & Justification


Layer	Choice	Why this, and what was rejected
Ingestion / messaging	Apache Kafka 4.1 (KRaft)	Replay, per-vehicle ordering through partition keys, consumer groups for scale-out. RabbitMQ has no cheap replay; Redis Streams are memory-bound; Pulsar has more moving parts (ADR-001)
Stream processing	Python consumer groups (confluent-kafka / librdkafka, orjson)	One language across the system; scale by partitions × processes; measured 25.8K msg/s per normalizer process before committing to it. Flink and Kafka Streams are JVM-heavy for an 8 GB laptop (ADR-005)
Batch processing	DuckDB over Parquet; ClickHouse SQL	Scans millions of rows on a laptop, no cluster. Spark rejected as too heavy
Relational	PostgreSQL 17	ACID, constraints for business rules, row-level security, good planner tooling (ADR-003)
Telemetry (NoSQL, columnar)	ClickHouse 25.8	Fast inserts and time-window scans, 4× compression measured, TTL tiers, native Kafka engine. TimescaleDB rejected for less ingest headroom on one node (ADR-002)
Cache / live state	Redis 7.4	Live state for 100K vehicles, sorted-set rankings, geo index, pub/sub fan-out, rate limits
Search / vector	none	Not needed: no free-text search, no embeddings
Backend	FastAPI	Async, OpenAPI generated, WebSockets
Frontend	React + TypeScript (Vite) , Leaflet	Typed UI, 133 KB gzipped, served by nginx on the API's origin
ML	LightGBM , TreeSHAP explanations	Tabular data, microsecond inference, built-in per-prediction contributions. Deep sequence models rejected: more data, harder to explain (ADR-007)
Infrastructure	Docker Compose (one command); Helm + Terraform planned (M16)	Runs on a laptop; the same images deploy to Kubernetes
CI/CD	GitHub Actions	Lint, types, unit, integration (Testcontainers), web, compose smoke with Playwright, coverage gate, 4 security scanners
Observability	Prometheus + Grafana , JSON logs	SLO burn-rate alerts first; distributed tracing deferred (ADR-009). ELK rejected for memory

Changed from the M0 plan, with reasons: MinIO was dropped (no longer published as community images, and costs RAM), and Parquet goes to `OBJECT_STORE_URL` instead. The canonical event is versioned JSON, not Protobuf, because the measured cost of JSON was acceptable and it

keeps every consumer simple (ADR-005). OpenTelemetry tracing is deferred (ADR-009). The radar uses exact counts, not a Count-Min Sketch (§3.3).

5.3 Data Architecture

ER diagram. The relational core has 20 tables in 3NF. The full diagram is in [er-diagram.md](#); the core entities are:



Deliberate denormalisation. `tenant_id` is repeated on every tenant-owned table, although on `vehicles` it follows from `fleet_id`. Row-level security and tenant-leading indexes need it without a join. The anomaly it could cause (a vehicle in another tenant's fleet) is blocked by composite foreign keys, and a test proves the database rejects it ([data model](#)). Business rules live in the schema: exclusion constraints, partial unique indexes, an append-only audit trigger, and a `CHECK` that every cost cites a source.

Polyglot map.

Data	Store	CAP choice	Retention
Tenants, vehicles, users, roles, costs	PostgreSQL	CP (single primary, synchronous commit)	Life of tenant
Alerts, work orders, audit log	PostgreSQL	CP	2 years (policy); audit by monthly partition
Raw canonical telemetry	ClickHouse <code>events</code>	AP (replays collapse by key)	30 days (hot)
Per-minute rollups	ClickHouse <code>vehicle_minute</code>	AP	400 days (warm)
Model scores, rejected payloads	ClickHouse	AP	180 / 7 days
Live state, rankings, rate limits	Redis	AP, rebuildable from Kafka	Ephemeral
Event log	Kafka	CP (acks=all ; min ISR 2 in the HA overlay)	1 day telemetry, 7 days alerts and DLQ
Training sets, cold archive	Parquet at <code>OBJECT_STORE_URL</code>	-	Per policy

At-least-once delivery plus idempotent writers joins the two sides. A reconciliation script proves no loss end to end ([ADR-004](#), [evidence](#)).

Capacity estimate (per-event sizes measured on simulated data; volumes are arithmetic):

Quantity	Value	Basis
Events per second	10K (100K vehicles every 10 s, load-tested); 100K (the brief)	-
Bytes per event	535 B raw JSON; 47.6 B on disk in ClickHouse (4.0× compression)	evidence , evidence

Quantity	Value	Basis
Events per day	0.86 billion (10K/s); 8.64 billion (100K/s)	× 86,400
Hot tier (raw, 30 days)	1.2 TB (10K/s); 12.3 TB (100K/s)	× 47.6 B × 30
Warm tier (rollups, 400 days)	2.4 TB	144M rows/day × 41 B × 400
Yearly raw volume if kept	15 TB (10K/s); 150 TB (100K/s)	× 365; this is why raw is kept only 30 days
Partition keys	Kafka: vehicle id; ClickHouse: day, sorted by (tenant, vehicle, time, seq); PostgreSQL audit: month	-


Query optimisation. The three slowest queries were found by replaying real dashboard traffic over 100K vehicles and a year of generated history (2M alerts), ranked with `pg_stat_statements` ([write-up](#), [evidence](#)):

Query	Before (ms)	After (ms)	Change made
Dashboard fleet summary	1,663 (mean)	0.97	8 correlated subqueries → one <code>FILTER</code> pass per table over active rows; partial indexes
Work orders by status, next page	58.2 (mean)	0.17	<code>ORDER BY</code> was bound to a text alias no index could serve → qualified columns; composite index
Planner input (active alerts without a work order)	309 (mean of 3 cycles)	473: not improved	A partial index helped when warm, not cold; the real fix is incremental planning (future work)

The whole replayed workload went from 195.6 s to 23.8 s of API time. Each rewrite is tested to return exactly what the straightforward query returns.

5.4 Deployment View

Today: one host, Docker Compose.



- **Replicas:** consumers scale with `NORMALIZER_REPLICAS` / `DETECTOR_REPLICAS`, up to the number of partitions (6 by default).
- **Networking:** every published port binds to `127.0.0.1` by default (`HOST_BIND`), and the dashboard proxies the API on the same origin.
- **Secrets:** `.env`, never committed. In production mode the API refuses to start with development defaults.
- **Limits:** every container has a memory limit; a unit test keeps the default total within 5 GB.
- **Autoscaling:** none locally.

Planned (M16): Kubernetes via Helm, cloud via Terraform. The same images, with configuration only: managed PostgreSQL and Redis with TLS and authentication, Kafka with SASL_SSL, consumers as Deployments scaled on consumer lag, TLS at the load balancer. **Not built yet.** This section will be updated with what is actually deployed.

Cloud-agnostic approach.

- Every dependency is an open protocol: the Kafka protocol, PostgreSQL wire, the Redis protocol, ClickHouse HTTP, and an S3-style object URL.
- Each can be self-hosted or replaced by any cloud's managed service by changing environment variables.
- The images do not change between the laptop and the cloud.


6. Low-Level Design

6.1 Layering & Separation of Concerns

Style: ports and adapters in the pipeline, a thin transaction-script API.

- **Pipeline.** The domain modules are pure Python with no I/O imports: `detector.py`, `planner.py`, `radar.py`, `dedup.py`, `adapters.py`, `processor.py`, `features.py`. They take dictionaries in and return records out, so they are unit tested without Kafka or databases (core coverage is 94–99%, [strategy](#)). The adapters around them do the I/O:
 - `kafka_loop.py`: consume, handle, produce, flush, then commit.
 - `sinks.py`: PostgreSQL and Redis writers.
 - `*_main.py`: wiring, signals, metrics.
- **API.** Each FastAPI router owns its SQL. With 17 endpoints, a repository layer would mostly forward calls. The price is that SQL sits in the presentation layer; this is recorded as debt (§12).

Layers of the detector service (dependencies point down only):



Layer	Responsibility	Must not	In Prognos
Presentation / API	HTTP, WebSocket, validation, auth checks, DTO mapping	Contain business rules or SQL	<code>routers/*</code> : validation, auth and DTOs; SQL is here too (debt)
Application / service	Use cases, transactions, orchestration	Depend on a specific database or broker	<code>detector_service.py</code> , <code>planner_main.py</code> (the planner transaction), <code>erasure.py</code>
Domain	Entities, value objects, business rules	Import framework or infrastructure code	<code>detector.py</code> , <code>planner.py</code> , <code>radar.py</code> , <code>dedup.py</code> , <code>adapters.py</code> : no I/O imports
Infrastructure	Repositories, broker clients, external APIs, caches	Leak vendor types into the domain	<code>kafka_loop.py</code> , <code>sinks.py</code> , <code>registry.py</code> , <code>deps.py</code>

Folder structure (two levels):

apps/	api/ · simulator/ · stream-processor/ · web/ · batch/
packages/	common/ (VIN, DTC, geohash, logs, roster) · schemas/ (canonical event JSON Schema)
ml/	src/prognos_ml/ · models/failure-7d-v1/ · tests/
database/	postgres/ (migrations, seeds) · telemetry/ (ClickHouse migrations)
kafka/	topics/ (topics.conf, create script)
infra/	docker/ · monitoring/ (prometheus, grafana) · helm/ · terraform/ · kubernetes/
tests/	unit/ · integration/ (+ features/) · contract/ · security/ · load/ · soak/ · chaos/
scripts/	seed, reconcile, load tests, SQL workload, charts, dashboards
docs/	product/ · architecture/adr/ · database/ · algorithms/ · api/ · security/ ·
	performance/ · observability/ · testing/ · ml/ · solution-document/
evidence/	benchmarks/ · load-tests/ · performance/ · chaos/ · security/ · coverage/ · sbom/

6.2 Design Principles Applied

SOLID:

- **Single responsibility.** One service per concern (normalizer, detector, sink, planner, radar, scorer), each deployable and scalable alone.
- **Open/closed.** A new OEM format is a new adapter function registered in `adapters.resolve()`; the rest of the pipeline does not change.
- **Liskov substitution.** Every OEM adapter returns the same `Canonical` record, so consumers can take any of them.
- **Interface segregation.** The loop exposes one small hook, `handle(msg, now)`.
- **Dependency inversion.** Domain code receives plain data. It never imports a client.

12-Factor App:

- **Config:** all of it comes from the environment (`.env.example`).
- **Stateless processes:** detector state is per partition and rebuilt on reassignment.
- **Disposability:** graceful shutdown flushes and commits.
- **Dev/prod parity:** the same images everywhere.
- **Logs:** JSON on stdout.
- **Admin processes:** one-off containers (`migrate`, `planner --once`).

Other principles:

- **Idempotency:**
 - Deterministic `event_ids` and alert fingerprints.
 - `ON CONFLICT DO NOTHING` upserts.
 - A partial unique index for active work orders.
 - `ReplacingMergeTree` in ClickHouse.
- **Fail fast:**
 - The API refuses unsafe production configuration.
 - The sink exits with a clear reason after its retry budget, rather than committing offsets it did not write.
- **Least privilege:**
 - API requests run as a restricted PostgreSQL role under RLS.
 - Column grants hide password hashes.
 - Containers run as non-root.
- **DRY:** one feature SQL for training and serving, proven identical by a parity test.
- **KISS:**
 - A greedy planner instead of an ILP.
 - Exact counts instead of a sketch.
 - JSON instead of Protobuf.

6.3 Design Patterns Used

Pattern	Problem it solves in Prognos	Location in code
Adapter	Maps each OEM's payload (flat metric, nested imperial, signal list) to one canonical event	apps/stream-processor/src/prognos_stream/adapters.py (orion_v1, vega_2_3, lyra_v1, resolve)
Retry with exponential backoff and full jitter, bounded by a time budget	A slow or restarting database must not lose alerts or crash-loop instantly	sinks.py (retry)
Idempotent consumer	At-least-once delivery without duplicate rows	sinks.py (ON CONFLICT), alert fingerprints in detector.py
Dead-letter queue	One bad message must not block a partition	processor.py → telemetry.dlq with a reason
Template method	One tested consume → produce → flush → commit loop shared by all consumers	kafka_loop.py (handle hook)
Strategy	Rank at-risk vehicles by rules or by the model	routers/vehicles.py (?source=rules or model)
Observer / publish-subscribe	Push alerts to every open dashboard without polling	Redis PUBLISH alerts:{tenant} in sinks.py → routers/live.py WebSocket
CQRS (light)	Writes go to PostgreSQL; the dashboard's live reads come from Redis projections	sinks.py LiveStateStore, routers/vehicles.py
Repository (partial)	Isolates alert persistence from the sink loop	sinks.py AlertStore
State machine	Work-order lifecycle (proposed → scheduled → in progress → completed/cancelled)	routers/work_orders.py + database CHECKs

Not used: an Outbox (Kafka is the source of truth and every writer is idempotent) and a circuit breaker (the API degrades live views when Redis is down and reports it in /readyz, but does not trip a breaker).

6.4 Interfaces, Contracts & Runtime Flows

API contract.


- **Documents:** [openapi.json](#), checked in and kept current by a test; the guide is [docs/api](#).
- **Versioning:** the URL path (/v1).
- **Pagination:** keyset (cursor, limit ≤ 200), stable while rows arrive.
- **Errors:** RFC 9457 application/problem+json with a request_id.
- **Rate limits:**
 - Per user (Redis, fixed window) for API calls.
 - Failed logins per address and account.
 - A 64 KB body limit at nginx.

Event schemas.


- **Streaming contract:** [asynccapi.yaml](#).
- **Canonical event:** JSON Schema [canonical-telemetry-v1](#), carrying schema_version.
- **Keys:** vehicle id on every per-vehicle topic, which keeps per-vehicle order.
- **Evolution:**
 - New optional fields are backwards compatible.
 - A breaking change is a new version number, handled by a new adapter.
 - Contract tests check each OEM format and the canonical schema (tests/contract).

- **Topics** ([topics.conf](#)): `telemetry.raw`, `telemetry.canonical`, `telemetry.dlq`, `alerts`, `vehicle.state` (compacted), `vehicle.risk` (compacted), `fleet.signals`, `sim.truth`.

Flow 1: a critical fault reaches the manager (measured at 100K vehicles).



Flow 2 (failure path): PostgreSQL goes away mid-ingestion (M13 drill, [write-up](#)).



A **duplicate event** takes the normal path and is dropped by the normalizer's sequence window (A6). A replay that slips past it collapses in ClickHouse and is ignored by the alert upsert. An **unknown OEM format** goes to `telemetry.dlq` with the reason `unknown_oem` or `unknown_schema_version`.

6.5 Algorithms & Data Structures

Full write-ups with pseudocode are in [algorithms.md](#).

Problem solved	Algorithm	Time / space	Scale tested (measured)
Fleet physics for 100K vehicles	Vectorised numpy state arrays, sharded across processes (A1)	$O(V)$ per tick, $O(V)$ memory	100K vehicles: 100K events/s on one core, 321K events/s on 4 (evidence)
Late and duplicate delivery	Min-heap of delayed messages (A3)	$O(\log B)$ per message	inside the runs above
Validation	ISO 3779 VIN check digit; SAE J2012 DTC parsing (A5)	$O(1)$	every event
Exact duplicate detection with out-of-order tolerance	IPsec-style anti-replay window: high-water mark + 1,024-bit bitmap per vehicle (A6)	$O(1)$ per event; $O(V \cdot W/8) \approx 13$ MB for 100K vehicles	1.19M messages, 0 duplicate event ids (evidence)
Early warning of degradation	Page's CUSUM, EW mean/variance, EWMA, time-weighted least-squares slope, with hysteresis (A7)	$O(1)$ per event per signal; 2.3 KB per vehicle	100K vehicles live (evidence)
Probability and deadline per rule	Right-censored frequencies with Wilson intervals; 10th-percentile time to failure (A8)	$O(n \log n)$	315 held-out alerts on another seed (evidence)
Workshop booking	Greedy by value, earliest deadline, over K nearest workshops $\times$ D days (A9)	$O(n \log n + n \cdot K \cdot D)$	5,000 alerts on 100K vehicles in 0.90 s (evidence)
Emerging faults	Distinct-vehicle counts per cohort, exact Poisson tail in log space, Bonferroni (A10)	$O(1)$ per DTC event; $O(\text{cohorts} \times \text{DTCs} \times \text{vehicles affected})$	695 ns per event, about 1.4M events/s (evidence)

Pseudocode of the duplicate window (A6), the hot path of every event:

```

check(v, seq):                                #  $O(1)$ 
  if no state[v]: state[v] = (seq, 1); return NEW
  high, bits = state[v]
  if seq > high: bits = (bits << (seq - high) | 1) & mask; high = seq; return NEW
  d = high - seq
  if d >= W: return T00_OLD                    # forwarded and flagged; ClickHouse is the backstop
  if bits >> d & 1: return DUPLICATE
  bits |= 1 << d; return NEW_LATE              # forwarded and flagged late

```

Pseudocode of the planner (A9):

```

candidates = one per (vehicle, failure mode): risk = max over its alerts,
  deadline = min(alert estimate, rule p10) - time since the alert
value(c)   = p × (unplanned - planned + downtime/day × extra days)  if costs are sourced
           = p × severity weight                                     otherwise
for c in sort(candidates, by value desc, deadline asc, vehicle id): #  $O(n \log n)$ 
  for day in horizon (D days), workshop in K nearest:                #  $O(K \cdot D)$ 
    if capacity(workshop, day) and day <= c.deadline: book; break
    else: book the earliest free slot and flag LATE, or leave unscheduled
insert all bookings in one transaction; ON CONFLICT DO NOTHING; audit each

```

7. Non-Functional Requirements & Performance Benchmarks

NFR	Target (case study)	Achieved	How measured
Ingest throughput	100K+ events/s	Not met end to end. Per stage: the simulator into Kafka held 100K events/s (evidence); the normalizer does 25.8K msg/s per process, 62.9K with 3 (benchmarks). The whole pipeline on 4 shared vCPUs sustains about 15K events/s and saturates at 20K (evidence). 100K vehicles reporting every 10 s (10K events/s) is sustained; the brief's rate of about one event per vehicle per second is not.	<code>scripts/load_test.py</code> sampling Prometheus
End-to-end latency	< 2 s dashboard; < 5 s critical alert	Met at 100K vehicles, 10K events/s: dashboard p95 1.88 s (96.7 % under 2 s); critical alert p95 1.78 s, p99 2.84 s	WebSocket client (evidence); detector histogram (evidence)
API latency	p95 < 200 ms; p99 < 500 ms	Met: p95 63-82 ms, p99 127-221 ms per route, 40 users during the 100K run, 0 errors	<code>scripts/api_load.py</code> , client-side (evidence)
Resilience	Recovers after broker / pod failure	Met in drills: detector stopped 13.5 min, then resumed from its committed offsets and drained a 7.4M-message backlog in 7 min 40 s (loss not separately reconciled in this drill; evidence); PostgreSQL stopped 90 s mid-ingestion, 0 loss over 3.77M events (evidence); one of 3 Kafka brokers killed: 1,000 of 1,000 <code>acks=all</code> writes succeeded, and replication healed after restart (evidence)	Drills with reconciliation
Availability	99.9 %, no single point of failure	Not met. Kafka can run 3 brokers (overlay), but PostgreSQL, ClickHouse and Redis are single instances locally. Availability over time is NOT YET MEASURED . The cloud design (M16) uses managed, replicated stores.	-

Load test setup.

- **Tools:** `scripts/load_test.py` seeds the fleet, drives the simulator and samples Prometheus and `docker stats`. `api_load.py` runs concurrent dashboard users, and `ws_latency.py` times alerts on the WebSocket.
- **Data volume:** 10K, 50K and 100K vehicles; a burst; a 30-minute soak; 50.7M events stored in ClickHouse over the runs.
- **Hardware:** one 4-vCPU x86_64 Linux container. The simulator, Kafka, every consumer and every database share those CPUs.


Results.

Scenario	Events/s in	Alert latency p50 / p95 / p99	CPU (of 4)	Memory	Verdict
10K vehicles	1,021	0.52 / 1.23 / 1.85 s	0.9	1.5 GiB	sustained
50K vehicles	5,015	0.55 / 1.57 / 1.92 s	1.9	2.4 GiB	sustained
100K vehicles	10,050	0.62 / 1.78 / 2.84 s	2.9	3.1 GiB	sustained
100K at 15K events/s	15,320	0.70 / 1.79 / 2.57 s	3.5	2.6 GiB	at the limit
100K at 20K events/s	20,322	18 / 29 / 30 s	3.8	3.5 GiB	not sustained
30-minute soak, 100K	10,106	0.59 / 1.83 / 5.04 s	2.9	3.0 GiB	sustained, memory flat

Source: [load tests](#). **Burst and soak graphs:**


A burst: the backlog peaks at 587K messages and drains 50 s after input returns to normal

100K vehicles at 10K events/s, 3x burst requested; the simulator delivered ~22K/s on the shared 4 vCPUs



30-minute soak: memory stays flat, the pipeline stays current

100K vehicles at 10K events/s



8. Security & Compliance

Threat model. A full STRIDE analysis covers the ingestion path and the public API. Each threat maps to its control, the code and a test ([threat model](#)). The top five:

#	Threat	Control	Test
1	Tenant A reads tenant B's vehicles (a missed filter)	Application filter and PostgreSQL row-level security under a restricted role; other tenants' ids answer 404 (ADR-010)	<code>test_row_level_security.py</code> runs unfiltered queries on purpose
2	Forged token (alg=none , HS256 key confusion, edited claims, wrong audience)	RS256 only, required claims, audience and issuer checked, 15-minute lifetime	<code>test_token_attacks.py</code> : 14 cases
3	A role calls an endpoint it should not	Permissions come from the server's <code>role_permissions</code> , never from token claims	<code>test_access_matrix.py</code> : 5 roles × 17 endpoints; a new route without a policy entry fails
4	Fake telemetry from a device	Schema validation and a registered-VIN check (unknown vehicles go to the DLQ). Residual: no per-device identity	normalizer tests
5	Databases reachable from the network	Ports bind to <code>127.0.0.1</code> (found and fixed in M12 : Redis and Kafka had been exposed on all interfaces)	<code>test_compose_security.py</code>

Authentication and authorisation.

- **Login:** OAuth2 password grant issuing RS256 JWTs, with the public keys at `/.well-known/jwks.json` (OIDC-ready). Passwords are hashed with `argon2id`, and a brute-force guard limits failed logins.
- **RBAC:** 5 roles, with permissions held server-side.
- **Tenant isolation:** in two layers (above).
- **Audit:** every change and every denial is written to an append-only log.

Device identity and encryption.

- **Not done locally:**
 - No per-device identity.
 - No TLS on internal links (all loopback or the Docker network).
 - No authentication on Redis or Kafka.
 - Encryption at rest is left to the disk.
- **Plan (M16):** an ingest gateway with per-device mTLS, Kafka SASL_SSL, managed stores with TLS and encryption at rest, and TLS at the load balancer.
- **Secrets:** in `.env` (git-ignored). Gitleaks runs on every push (0 findings, [evidence](#)), and production mode rejects default secrets.

Privacy ([privacy](#)).

- **Minimisation:** drivers exist only as tenant-issued pseudonyms (`DRV-123456`); no names or licence numbers are stored.
- **Masking:** analysts see positions rounded to about 1 km.
- **Retention:** location history is kept 30 days by default, per tenant.
- **Erasure:**
 - A data protection officer files a request.

- A worker removes the vehicle's location history from ClickHouse, Redis and the driver links, auditing each step and recording what remains and until when (Kafka retention).
- Live run: 1,935 location rows erased in about 14 s in a 19.2M-row table ([evidence](#)).
- **Law:** designed against GDPR (Art. 5 minimisation, Art. 17 erasure) and India's DPDP Act 2023. No legal review was done.

AI safety. There is no LLM or agent, so prompt injection and tool permissions do not apply. The ML model is guarded instead:

- **Shadow mode:** it cannot drive the planner until validated ([ADR-008](#)).
- **Data gate:** it scores only vehicles with enough recent data.
- **Human approval:** every work order is a **proposal** that a person schedules.
- **Audit:** proposals are audited with their reasons.

Supply chain.

- **Trivy** fails CI on any HIGH or CRITICAL finding in dependencies (0 today, [evidence](#)) and on fixable CRITICAL findings in built images (0; HIGH findings in base-image OS packages are reported, [evidence](#)).
- **Semgrep** runs on every push.
- **Dependabot** groups minor and patch updates weekly.

9. Test Strategy

Test type	Tools	No. of tests	Coverage / result	In CI?
Unit (with contract and security)	pytest	233, plus 33 documentation checks (links, anchors, traceability)	84 % branch coverage, unit + integration combined, all 5 Python packages; gate at 80 % (report)	Yes
Integration & contract	pytest + Testcontainers (real PostgreSQL 17, ClickHouse 25.8, Kafka, Redis); JSON Schema contracts	80 (includes the 5 BDD scenarios)	all pass	Yes
Acceptance (BDD)	pytest-bdd (Gherkin)	5 scenarios	5 / 5 pass: alert → scheduled repair, idempotent planning, tenant isolation, role limits, location masking	Yes
Web	Vitest; Playwright against the nginx container	7 + 4	all pass; includes security headers	Yes
Monitoring rules	promtool	11 rule tests	each alert fires on a synthetic failure and stays quiet on normal traffic	Yes
Performance / load / soak	<code>load_test.py</code> , <code>api_load.py</code> , <code>ws_latency.py</code>	10K / 50K / 100K, 15K and 20K events/s, burst, 30-min soak, 40-user API run, 2 freshness runs	100K vehicles at 10K events/s sustained; alerts p95 1.78 s; API p95 < 82 ms; dashboard p95 1.88 s	No (recorded evidence)
Security (SAST, dependency, image, secrets)	Semgrep, Trivy (fs + image), gitleaks	4 scanners	0 HIGH/CRITICAL dependency findings; 0 fixable CRITICAL in images; 0 secrets. No DAST (not run)	Yes
Compliance & chaos	pytest (RLS, erasure, audit); manual drills	3 drills (Kafka broker, detector, PostgreSQL)	Kafka: 1,000/1,000 writes with a broker down; PostgreSQL: 0 loss over 3.77M events by reconciliation; detector: recovered from committed offsets	Tests yes; drills recorded

Edge cases covered.

- **Delivery:** duplicates (1 % injected), out-of-order events (2 %, up to 30 s late), malformed payloads (7 kinds), missing fields, unknown OEM format and unknown vehicle (all to the DLQ with a reason), clock skew, and a simulator restart mid-stream.
- **Outages:** broker, detector and database outages.
- **Isolation and access:** cross-tenant ids, forged tokens, and a role calling a forbidden endpoint.

Sample input / output. A VEGA payload (nested, imperial, IST timestamp, DTCs as one string) becomes a canonical event with metric units, UTC and a DTC list; see `tests/contract/test_oem_contract.py` and [the schema](#).

Can the tests fail? Several were checked by breaking the code on purpose:

- Disabling location masking fails the analyst scenario.
- Serving the old nginx configuration fails the headers test.
- Each login-concurrency test fails on the version it guards against.

Details: [strategy](#).

Reports: [coverage](#) · [load tests](#) · [Trivy](#) · [image scan](#) · [gitleaks](#).

10. Observability

Dashboard screenshot. The service-health dashboard (generated from code) during the detector outage drill shows throughput, seconds behind per stage, alert latency percentiles and API errors:



SLOs (definitions):

SLO	Target
Critical alerts raised within 5 s	95 %
API availability	99.5 %
Data freshness	< 60 s behind
Data quality	< 5 % to the DLQ

Each SLO is a Prometheus recording rule with multi-window burn-rate alerts, and every alert links to a [runbook](#).

Troubleshooting walk-through: a latency spike. This is what happened in M14 at 100K vehicles:

1. **The symptom pages.** `AlertLatencySLOBurn` fires: more than 5 % of critical alerts take longer than 5 s.
2. **Find the stage.** On the dashboard, the "seconds behind the stream" panel per consumer group shows which stage is late. Here the detector was 150 s behind, handling about 2.2K of 10K messages per second.
3. **Healthy or restarting?** The target answered scrapes between restarts, so `ServiceDown` stayed quiet, but the detector's process start time kept changing: it was crash-looping (25 restarts in 5 minutes). That gap led to a new alert, `ServiceRestarting` (`changes(process_start_time_seconds[15m]) > 2`).
4. **Why.** `docker stats` and the kernel log showed the process OOM-killed at its 512 MB limit (`Memory cgroup out of memory`). The memory benchmark found that each vehicle's state held a whole event. The fix made per-vehicle state 54 % smaller.
5. **Confirm.** The rerun shows 0 restarts and p95 back to 1.78 s. The JSON logs carry a `request_id` (API) and `event_id` (pipeline), which correlate a single request or event across services.

Distributed tracing was **deferred** ([ADR-009](#)): for a pipeline whose stages are separated by Kafka, per-stage lag and latency histograms answer "where is the time going" more cheaply than traces.

11. AI / ML Component

Purpose. Rank vehicles by the probability of a breakdown within 7 days, so the workshop's limited capacity goes to the right vehicles. Rules alone fire late and treat every firing alike. In the held-out back-test, the rules' top-5 % list was 72.5 % precise; the model's was 99.7 %.

Data and features.

- **Data:** five simulated 8-hour runs of 300 vehicles (about 8.5M events each), passed through the production normalizer and detector and stored as Parquet ([evidence](#)).
- **Features:** DuckDB computes per-minute buckets, then 60-minute and 10-minute window statistics: DTC rates, minimum 12 V voltage, the tyre pressure ratio, peak coolant.
- **Leakage checks:** a test proves no feature looks past the prediction time. Train and test never share a simulated run, so test vehicles, faults and noise are unseen; early stopping validates on training vehicles held out from fitting.
- **Serving parity:** serving computes the same SQL on ClickHouse data, and a parity test proves identical features and predictions.

Model design.

- **Model:** LightGBM, a binary objective, 94 trees.
- **Artefact:** versioned, with a SHA-256 checksum. Loading refuses a mismatched file or feature list.
- **Explanations:** TreeSHAP for every prediction, shown to the manager in plain words.
- **Out of scope:** no vector store, no agent, no tools.

Evaluation against the baseline the planner uses today (calibrated rules), on the same rows ([model card](#), [evidence](#)):

Test set	Model PR-AUC	Rules PR-AUC	Gain (95 % CI)
Held-out run	0.742	0.508	+0.230 [0.171, 0.302]
Slower faults	0.785	0.645	+0.139 [0.101, 0.180]
Rarer faults	0.702	0.550	+0.147 [0.057, 0.236]
v3, vehicles reporting every 10 s	0.679	0.478	+0.20 [0.15, 0.25] (evidence)

Guardrails and cost.

- **Shadow mode:** the first live run put 8 % of the fleet above 0.5, because one feature depended on the reporting rate. The model was fixed (v3), put behind a data gate (at least 30 events covering 45 of the last 60 minutes), and kept in shadow; the planner still uses the rules ([ADR-008](#)).
- **Human in the loop:** work orders are proposals.
- **Hallucination:** not applicable (no generative model).
- **Cost:** 0.024 ms per vehicle (p50); a live scoring cycle over 100K vehicles takes about 17 s on 4 vCPUs, mostly feature SQL ([evidence](#)). No per-request API cost.

Limits (read before quoting any number).

- The data is simulated, so the model may be learning the simulator.
- The median lead of 50 h misses the 72 h target.
- Recall at ≥ 48 h is 35 %.
- Money is not computed, because costs are not sourced.

12. Architecture Decisions, Risks & Future Enhancements

Architecture decision records (all ten are in [docs/architecture/adr](#)). The five that shape the system:

ADR	Context	Options considered	Decision	Consequences
001 Messaging	100K events/s, per-vehicle order, replay after a bug	Kafka, RabbitMQ, Pulsar, Redis Streams	Kafka (KRaft)	Replay and consumer-group scaling; one more stateful system to run
002 Telemetry database	Billions of rows, time-window scans	ClickHouse, TimescaleDB, Cassandra	ClickHouse	47.6 B/event on disk and a native Kafka engine; eventual consistency and merge-time dedup
003 Relational core	Money, work orders, audit must be exact	PostgreSQL, MySQL	PostgreSQL	Constraints and RLS enforce rules; a single primary to scale up, not out
004 CAP trade-off	Different data needs different guarantees	One store for all; per-class choice	CP for alerts, work orders, Kafka writes; AP for telemetry history and live state, made safe by idempotent keys	A partitioned store may refuse writes rather than lose them; live views may be briefly stale; a reconciliation job proves no loss
008 ML rollout	The model beat the rules offline but misbehaved on live data	Switch over; shadow; drop	Shadow mode with a data gate and rate-invariant features	No unvalidated model drives bookings; live outcomes are needed to switch

Risks and technical debt.

1. **Throughput.** End to end, 15K events/s on 4 vCPUs, against the brief's 100K+. The stages scale by partition, but scaling out on real cores is **NOT YET MEASURED**.
2. **Simulated world.** Every quality number comes from our own simulator, and real-fleet performance is unknown.
3. **Money not sourced.** Cost parameters are placeholders, so no saving is claimed.
4. **Single-node stores and no transport security** locally (Redis and Kafka without authentication or TLS; no device identity).
5. **SQL in API routers** (§6.1). Also the planner reads all active alerts each cycle, and M15 showed that query did not improve.
6. **Redis 7.4 is source-available** (RSALv2/SSPLv1), not open source ([open-source list](#)). Valkey is a drop-in candidate, not yet tested.
7. **Not measured on the target Mac**; the M0 plan's 300K-event burst was not reached (about 22K/s delivered on shared CPUs).
8. **Fixable HIGH CVEs** in base-image OS packages (reported, not gating).

Future enhancements: three steps from prototype to pilot.

1. **Deploy securely (M16).**
 - Helm and Terraform on one cloud.
 - Managed PostgreSQL and Redis with TLS.
 - Kafka with SASL_SSL and an ingest gateway with per-device mTLS.
 - Consumers autoscaled on lag.
 - Then re-measure throughput on dedicated cores.
2. **Pilot on real data.**
 - One real fleet's telemetry through a new OEM adapter.

- Sourced repair and downtime costs (turning on money).
- Work-order outcomes fed back to recalibrate the rules and validate the model on live outcomes, the condition for leaving shadow mode.

3. **Scale the hot path.**

- Incremental planning (only new or changed alerts).
- More partitions.
- If Python per-core throughput is the limit on real cores, port the normalizer to Go behind the same topics and contracts (the ADR-005 fallback).

13. Demo Video (5 Minutes Maximum)

Time	Segment	What to show
0:00 - 0:30	Problem	The fleet manager, the breakdown problem, and one measured number
0:30 - 1:00	Solution	One-line pitch and the value proposition
1:00 - 3:00	Live demo	A scripted failure on live data: the critical alert arrives on the dashboard within seconds; acknowledge; the planner proposes a booking; schedule it; the vehicle detail explains why
3:00 - 4:15	Under the hood	Architecture; the 100K-vehicle load test and Grafana; one failure-recovery moment (a stopped service and its backlog draining)
4:15 - 5:00	Impact & next steps	Measured results, what is not yet proven, next steps, the team

Script: <docs/demo/demo-script.md>, timed from a real rehearsal ([timeline](#), [recovery moment](#)).

Video: in folder **2 - Hackathon Explainer Video** of the submission Google Drive folder.

14. Repository Checklist

Item	Status	Where
README: problem, architecture, quick start, environment variables, test commands, known issues	Done	README
One-command run	Done: <code>make up</code> starts the data plane; <code>make pipeline-demo</code> adds the simulator, pipeline, API and dashboard	Makefile
Structure: service folders, <code>/docs</code> , <code>/infra</code> , <code>/tests</code>	Done	§6.1
CI: build, lint, all test suites and security scans on every push	Done: 8 jobs across <code>ci.yml</code> and <code>security.yml</code>	.github/workflows
No secrets committed	Done: gitleaks on every push, 0 findings	evidence
<code>.env.example</code> provided	Done	.env.example
Commits from all members	Solo project: all commits by the author, assisted by Claude Code (§16)	git history
Final tag <code>v1.0-submission</code>	Done at submission	git tags

15. Conclusion

Key learnings.

- **Measurement finds what review does not.** Load tests found three production-grade bugs: a memory crash loop at 100K vehicles, an API that froze during a login burst, and a Kafka setting that filled the disk. The SQL workload found an `ORDER BY` that no index could serve. The failure drills found alerts that could not see a service vanishing from discovery.
- **A model that wins offline can fail live.** The model's first live run flagged 8 % of the fleet because one feature depended on the reporting rate. Shadow mode caught this before it drove any booking.
- **Honesty is a design constraint.** "Not sourced" and "not yet measured" are shown in the product and the documents, and a database constraint enforces them for costs.

Strengths.

- End to end on one machine, with measured latency inside the brief's targets at 100K vehicles.
- Correctness proven by reconciliation (no loss, no duplicates) and by tests that were checked to fail.
- Two layers of tenant isolation.
- Every claim linked to evidence.

Challenges tackled.

- **Per-core throughput.** Python at 100K events/s: measured per stage first (M4), then scaled by partition.
- **Honest evaluation.** No labelled public data, so the simulator supplies ground truth.
- **An 8 GB laptop.** Solved with memory limits, an opt-in observability profile and an HA overlay only for drills.

16. Declarations

Open-source components.

- **The full list with licences:** <docs/open-source.md>, generated from CycloneDX SBOMs of the built images in <evidence/sbom/>.
- **Main components:** Kafka, ClickHouse, Prometheus and Trivy are Apache-2.0; PostgreSQL is under the PostgreSQL License; FastAPI, LightGBM, DuckDB, React and dbmate are MIT; Leaflet is BSD-2-Clause; Grafana is AGPL-3.0; psycopg is LGPL-3.0.
- **Redis 7.4** is source-available (RSALv2/SSPLv1), not open source; it is used self-hosted, which its licence allows.
- **Prognos itself:** MIT ([LICENSE](#)).

AI tools used.

- **Claude Code (Anthropic)** was the coding assistant throughout: design discussion, code, tests, measurement scripts and documentation.
- Every result was verified by running it. Numbers come from evidence files produced by scripts in the repository, not from the assistant.
- The product itself uses no LLM or external AI API.

Data. All data is synthetic, generated by this repository's simulator. It contains no real personal or vehicle-owner data. Drivers are pseudonyms, and VINs belong to fictitious generated vehicles.

Affiliation. Motorq is referred to only as an industry example. This is an independent academic project with no affiliation to Motorq.

17. Appendix

Documents.

Topic	Document
Problem selection, decision matrix, MVP	m0-problem-selection.md
Architecture decisions	ADR-001 ... ADR-010, CAP / PACELC
Data model and ER diagram	data-model.md , er-diagram.md
Algorithms with pseudocode	algorithms.md
API	README , OpenAPI , AsyncAPI
Detection baseline	detection-baseline.md
Benchmarks, load tests, SQL	benchmarks.md , load-tests.md , sql-optimisation.md
Model card	model-card.md
SLOs, runbooks, drills	slo.md , runbooks.md , drills.md
Security and privacy	threat-model.md , privacy.md
Testing	strategy.md
Feature traceability	feature-traceability.csv
Final audit: every requirement → code → test → evidence → demo	final-audit.md
Demo script	demo-script.md
Open-source components	open-source.md

Evidence index.

- [evidence/benchmarks/](#) : stage benchmarks, back-tests, model results, storage.
- [evidence/load-tests/](#) : M14 scenarios and M17 dashboard freshness.
- [evidence/performance/](#) : SQL plans before and after.
- [evidence/chaos/](#) : failure drills.
- [evidence/security/](#) : scan reports, erasure run.
- [evidence/coverage/](#) : coverage report.
- [evidence/sbom/](#) : SBOMs.

References.

- Beyer et al., *The Site Reliability Workbook*, "Alerting on SLOs" (O'Reilly, 2018): burn-rate alerts.
- E. S. Page, "Continuous inspection schemes", *Biometrika* 41 (1954): CUSUM.
- E. B. Wilson, "Probable inference, the law of succession, and statistical inference", *JASA* 22 (1927): the Wilson interval.
- S. Kent, RFC 4303, *IP Encapsulating Security Payload* (2005), §3.4.3: the anti-replay window.
- M. Nottingham et al., RFC 9457, *Problem Details for HTTP APIs* (2023).
- M. Brooker, "Exponential Backoff and Jitter", AWS Architecture Blog (2015).
- Lundberg et al., "From local explanations to global understanding with explainable AI for trees", *Nature Machine Intelligence* 2 (2020): TreeSHAP.
- ATRI, *An Analysis of the Operational Costs of Trucking* (2026 update), via the trade press cited in §2.2.
- ISO 3779 (VIN); SAE J2012 (diagnostic trouble codes).